

Lecture V: The Item RAM Strikes Back!

Towards a practical dynamic bulk storage system for use in the real world

Instructor: Andrews54757

CSE269: Introduction to Encoded Storage

S∞ntech Annals

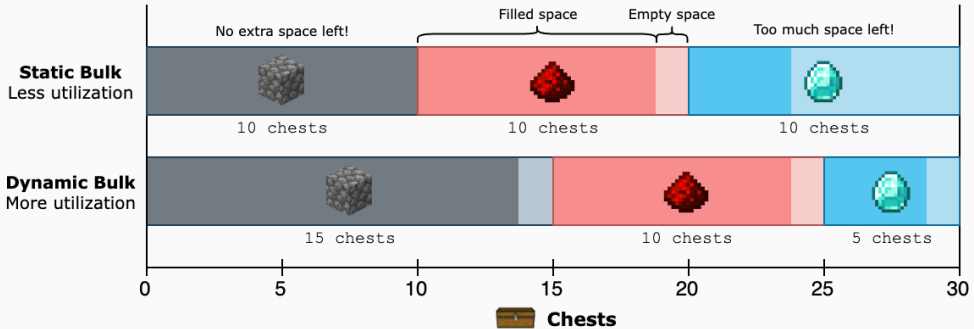
Overview

1. What is dynamic bulk?
2. Fast and compact data storage is hard
 - 2.1 Internal memory is too big, complex, and laggy
 - 2.2 Item RAMs are still too big
 - 2.3 Disk Drives are too slow
3. Solution: Linked-list based dynamic bulk
 - 3.1 Non-box item RAM is fast and compact
 - 3.2 Linked-list to the rescue
4. Conclusion and future steps




1. What is dynamic bulk?

- **Static bulk** has a fixed capacity for each item type
- **Dynamic bulk** automatically expands and shrinks capacity for items




1. What is dynamic bulk? Contd.

- Divide your storage into **slices** of fixed size
- Each slice can store a fixed number of items
- The size of your slice determines the minimum non-zero amount of storage space you can allocate to an item type




1. What is dynamic bulk? Contd.

- Let's store 8 chests of redstone
 - Do we have a slice assigned to redstone? **No!**
 - Allocate an available slice to redstone
 - Now we have space to start storing redstone!




1. What is dynamic bulk? Contd.

- Let's store 8 chests of redstone
 - Fully filled one chest with 5 chests of redstone
 - But we have more 3 more chests of redstone to store!




1. What is dynamic bulk? Contd.

- Let's store 8 chests of redstone
 - One slice is **fully** filled with redstone (5 chests)
 - Another slice is **partially** filled with redstone (3 chests / 5 chests)






1. What is dynamic bulk? Contd.

- Retrieve 1 chest of redstone
 - Pull from the partially filled slice first!
 - Minimize internal fragmentation



2. Fast and compact data storage is hard

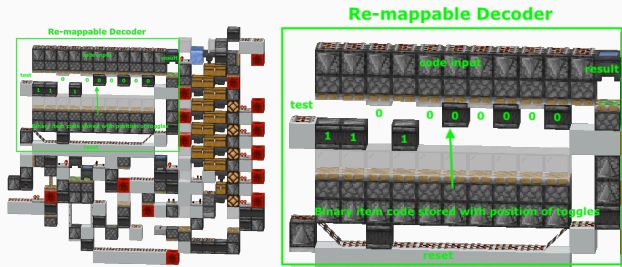
- We need to keep track of what slices are assigned to each item type
- Must store that information in a way that is quickly accessible
- Storage must be dense so that it is compact



From left to right: PallaPalla's Dynamic Slice, 1000 Item RAM, Obi's Disk Drive

2.1 Internal memory is too big, complex, and laggy

- PallaPalla, original inventor of dynamic bulk, uses a binary re-mappable decoder to store assigned item code within each slice.
- Almost every piston in each slice's decoder is fired to write/test item codes.
- Allocation logic is internal to the slice, so each slice is large and complex.



2.1 Internal memory is too big, complex, and laggy. Contd.

- What if we externalized the slice allocation logic?
 - Each slice is now just a simple static encoded bulk unit
 - Store information about slice allocation in a separate memory unit
 - Externalized logic is simpler and more compact




Figure 12: Diagram of dynamic bulk with externalized logic

2.1 Internal memory is too big, complex, and laggy. Contd.

- How is slice allocation information stored externally?
 - Represent slices with an encodeable item called the **slice code item**. One slice code item corresponds to one slice.
 - Maintain a list of slice code items in a box dedicated to one item type
 - Store the list of slice code items in an external item memory unit in the address corresponding to the item type




Figure 12: Externalized dynamic bulk data layout

2.2 Item RAMs are still too big

- An item RAM is a memory unit that stores items in an addressable way
- Each item type is assigned a unique address
- The item RAM can be read from and written quickly in constant time ($O(1)$)




Figure 13-14: 1000 address item RAM and its slice. 46gt call latency.

2.2 Item RAMs are still too big. Contd.

- Unfortunately, item RAMs are very big.
- Latest advancements have doubled the density, but still too big.
- 17 x 41 x 33 blocks for 1000 addresses



Figure 15-16: Improved 1000 address item RAM and its slice. 46gt call latency.

2.3 Disk Drives are too slow

- What if, we stored multiple boxes from different addresses in a single block?
- Store up to 54 boxes in a double chest
- Cycle through the boxes to find the correct one



Figure 17-18: Obi's Disk Drive, 256 addresses.

2.3 Disk Drives are too slow. Contd.

- Cycling through slots is slow, 8gt per slot.
- Box at the 32nd slot takes a minimum 248gt to access, which is 12.4 seconds.
- Need something faster!



Figure 19: Minimum time in seconds to reach slot in a disk drive

3. Towards a better dynamic bulk

- How can we store the slice allocation information in a way that is fast and compact?
- Need a data storage that uses minimal space and can be accessed quickly
- What if we separate the partial slice allocation information from the list of full slices?

3.1 Non-box item RAM is fast and compact

- If instead of storing boxes, we store a single non-box slice code item per address
- Can pack multiple addresses in a single shulker box for higher density



Figure 20: Non-box item RAM storage scheme

3.1 Non-box item RAM is fast and compact. Contd.

- We can quickly bring boxes to an ultrafast slot cyclor to access the correct item
- Uses hoppercars, which can pull items from shulker boxes every game tick
- 8x faster cycling than traditional disk drives. Item at the 20th slot takes only 1 second



Figure 21-22: 1 slot/gt slot cyclor with internal layout

3.1 Non-box item RAM is fast and compact. Contd.

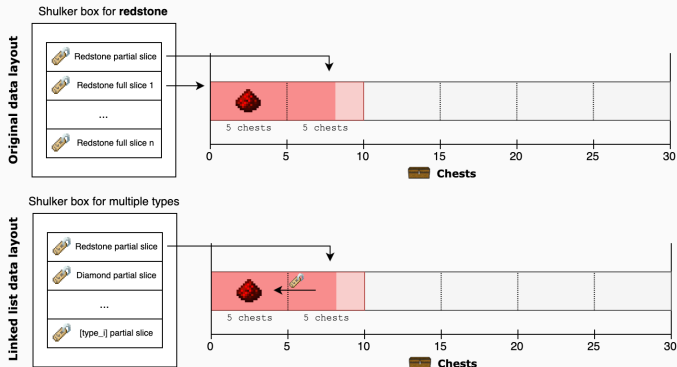
- Combine slot cycler with 50 item RAM (21gt call latency)
- Can access 1000 addresses within 55gt
- Result: Both fast and compact!



Figure 23: 1000 address non box item RAM

3.2 Linked-list to the rescue

- **Recall:** We need to be able to allocate more than one slice to an item type, but the non-box item RAM can only store one slice per address
- **Solution:** Store the next full slice's slice code item in the current slice to form a linked list



3.2 Linked-list to the rescue. Contd.

- Maintaining the linked list for **insertion** is simple.
 1. Obtain a slice code item at a slot corresponding to the item code using the 1000 non-box-item RAM.
 2. Insert boxes into the corresponding slice.
 3. If slice is full, pull a new slice code item from storage to allocate a new empty slice
 4. Put the now-full slice code item into the newly allocated slice and put the rest of the boxes into the newly allocated slice. Repeat 3-4 as needed.
 5. Store the remaining slice code item into the 1000 non-box-item RAM at the slot corresponding to the item code to finish insertion.

3.2 Linked-list to the rescue. Contd.

- Maintaining the linked list for **retrieval** is also simple.
 1. Obtain a slice code item at a slot corresponding to the item code using the 1000 non-box-item RAM.
 2. Retrieve boxes from the corresponding slice. Test items if they are boxes, if one is not a box then hold it for later, it is the slice code item for the next slice.
 3. If slice is empty and you have found a code item for the next slice, then switch to the next slice. Put the now empty slice code item back into storage.
 4. Keep retrieving boxes, repeat 2-3 as needed.
 5. If you have found the slice code item for the next slice, put it back into the slice.
 6. Store the remaining slice code item into the 1000 non-box-item RAM at the slot corresponding to the item code to finish insertion.



4. Conclusion and future steps

- Linked-list based dynamic bulk has the potential to be both fast and compact
- Future work: Implement the linked-list based dynamic bulk and test its performance
- Further optimizations: Reduce the latency of the non-box item RAM
- **This is an assignment for you!** Unfortunately I have no time to implement this myself. Good luck!
 - I have posted the schematic of the 1000 address non-box item RAM in the description to help you get started

